

fraud-detection-20188

May 24, 2024

#####Proyecto 2 - Security Data Science##### # #####Modelos de Deteccion de Fraude Bancario#####

0.1 Nombre: Pedro Arriola

0.2 Carnet: 20188

Este trabajo práctico tiene como objetivo investigar la viabilidad del entrenamiento incremental en modelos de aprendizaje automático y profundo, utilizando como estudio de caso un dataset de transacciones de tarjeta de crédito clasificadas en normales y fraudulentas. Los modelos a investigar den incluir 2 de los siguientes algoritmos: Redes Neuronales Artificiales (ANN), LightGBM, XGBoost, Random Forest y Máquinas de Vectores de Soporte (SVM).

0.2.1 Carga del Dataset

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Cargar el dataset
file_path = '/content/drive/MyDrive/fraud/fraud_feature_engineering_example.csv'
df = pd.read_csv(file_path)

# Exploración inicial del dataset
df.head()
df.info()
df.describe()


# Verificar valores nulos
df.isnull().sum()

# Convertir 'trans_date_trans_time' a datetime
df['trans_date_trans_time'] = pd.to_datetime(df['trans_date_trans_time'])

# Seleccionar solo columnas numéricas para la matriz de correlación
numerical_columns = df.select_dtypes(include=[float, int]).columns
corr_matrix = df[numerical_columns].corr()
plt.figure(figsize=(15, 10))
sns.heatmap(corr_matrix, annot=True, fmt='.2f')
```

```
plt.title('Matriz de Correlación')
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1852394 entries, 0 to 1852393
Data columns (total 27 columns):
 #   Column                                Dtype
---  -
 0   trans_date_trans_time                object
 1   cc_num                              int64
 2   merchant                            object
 3   category                            object
 4   amt                                 float64
 5   first                               object
 6   last                                object
 7   gender                              object
 8   street                              object
 9   city                                object
10   state                               object
11   zip                                  int64
12   lat                                 float64
13   long                                float64
14   city_pop                            int64
15   job                                 object
16   dob                                 object
17   trans_num                           object
18   unix_time                           int64
19   merch_lat                           float64
20   merch_long                           float64
21   is_fraud                            int64
22   amt_month                           float64
23   amt_year                            float64
24   amt_month_shopping_net_spend        float64
25   count_month_shopping_net            float64
26   first_time_at_merchant              bool
dtypes: bool(1), float64(9), int64(5), object(12)
memory usage: 369.2+ MB
```



0.2.2 Ingeniería de Variables

```
[ ]: # Crear variables temporales
df['Transaction_hour'] = df['trans_date_trans_time'].dt.hour
df['Transaction_day'] = df['trans_date_trans_time'].dt.day
df['Transaction_month'] = df['trans_date_trans_time'].dt.month
df['Transaction_year'] = df['trans_date_trans_time'].dt.year
df['Transaction_day_of_week'] = df['trans_date_trans_time'].dt.dayofweek
df['Transaction_weekend'] = df['Transaction_day_of_week'].isin([5, 6]).
    ↳astype(int)
df['Transaction_quarter'] = df['trans_date_trans_time'].dt.quarter

# Crear variables de días desde la última y primera transacción del mismo
    ↳cliente
df = df.sort_values(by=['cc_num', 'trans_date_trans_time'])
df['Days_since_last_transaction'] = df.
    ↳groupby('cc_num')['trans_date_trans_time'].diff().dt.days
df['Days_since_last_transaction'].fillna(0, inplace=True) # Rellenar NaN con 0
    ↳para la primera transacción
```

```

df['Days_since_first_transaction'] = df['trans_date_trans_time'] - df.
    ↳groupby('cc_num')['trans_date_trans_time'].transform('min')
df['Days_since_first_transaction'] = df['Days_since_first_transaction'].dt.days

# Verificación de las nuevas características
df_temporal_features = df[['Transaction_hour', 'Transaction_day',
    ↳'Transaction_month', 'Transaction_day_of_week',
    ↳'Transaction_weekend', 'Transaction_quarter',
    ↳'Days_since_last_transaction', 'Days_since_first_transaction']]
df_temporal_features.head()

```

```

[ ]:
Transaction_hour  Transaction_day  Transaction_month  \
1017              12              1              1
2724              8              2              1
2726              8              2              1
2882              12             2              1
2907              13             2              1

Transaction_day_of_week  Transaction_weekend  Transaction_quarter  \
1017                    1                    0                    1
2724                    2                    0                    1
2726                    2                    0                    1
2882                    2                    0                    1
2907                    2                    0                    1

Days_since_last_transaction  Days_since_first_transaction
1017                       0.0                          0
2724                       0.0                          0
2726                       0.0                          0
2882                       0.0                          0
2907                       0.0                          1

```

```

[ ]: # Creación de variables basadas en el monto
df['Transaction_amount_log'] = np.log1p(df['amt'])
df['Transaction_amount_square'] = df['amt'] ** 2
df['Transaction_amount_sqrt'] = np.sqrt(df['amt'])
df['Transaction_amount_cubed'] = df['amt'] ** 3

# Variables de estadísticas por cliente
df['Avg_transaction_amount'] = df.groupby('cc_num')['amt'].transform('mean')
df['Std_transaction_amount'] = df.groupby('cc_num')['amt'].transform('std').
    ↳fillna(0)
df['Min_transaction_amount'] = df.groupby('cc_num')['amt'].transform('min')
df['Max_transaction_amount'] = df.groupby('cc_num')['amt'].transform('max')

# Verificación de las nuevas características

```

```
df_monto_features = df[['Transaction_amount_log', 'Transaction_amount_square',
↳ 'Transaction_amount_sqrt',
                        'Transaction_amount_cubed', 'Avg_transaction_amount',
↳ 'Std_transaction_amount',
                        'Min_transaction_amount', 'Max_transaction_amount']]
df_monto_features.head()
```

```
[ ]:      Transaction_amount_log  Transaction_amount_square \
1017                2.112635                52.8529
2724                3.987872                2802.6436
2726                4.419804                6737.1264
2882                3.577669                1210.3441
2907                3.338613                738.7524

      Transaction_amount_sqrt  Transaction_amount_cubed \
1017                2.696294                384.240583
2724                7.275988                148371.952184
2726                9.059801                552983.334912
2882                5.898305                42107.871239
2907                5.213444                20079.290232

      Avg_transaction_amount  Std_transaction_amount  Min_transaction_amount \
1017                59.257796                142.869746                1.02
2724                59.257796                142.869746                1.02
2726                59.257796                142.869746                1.02
2882                59.257796                142.869746                1.02
2907                59.257796                142.869746                1.02

      Max_transaction_amount
1017                3437.46
2724                3437.46
2726                3437.46
2882                3437.46
2907                3437.46
```

```
[ ]: # Variables de frecuencia
df['Num_transactions_last_day'] = df.groupby('cc_num')['trans_date_trans_time'].
↳ transform(lambda x: (x.diff().dt.days.fillna(0) == 0).cumsum())
df['Num_transactions_last_week'] = df.
↳ groupby('cc_num')['trans_date_trans_time'].transform(lambda x: (x.diff().dt.
↳ days.fillna(0) < 7).cumsum())
df['Num_transactions_last_month'] = df.
↳ groupby('cc_num')['trans_date_trans_time'].transform(lambda x: (x.diff().dt.
↳ days.fillna(0) < 30).cumsum())
```

```

df['Avg_transactions_per_day'] = df.groupby(['cc_num',
↳df['trans_date_trans_time'].dt.date])['trans_date_trans_time'].
↳transform('count') / df.groupby('cc_num')['trans_date_trans_time'].
↳transform(lambda x: x.dt.date.nunique())

# Verificación de las nuevas características
df_frecuencia_features = df[['Num_transactions_last_day',
↳'Num_transactions_last_week', 'Num_transactions_last_month',
↳'Avg_transactions_per_day']]
df_frecuencia_features.head()

```

```

[ ]:      Num_transactions_last_day  Num_transactions_last_week  \
1017                                1                        1
2724                                2                        2
2726                                3                        3
2882                                4                        4
2907                                5                        5

      Num_transactions_last_month  Avg_transactions_per_day
1017                                1          0.001464
2724                                2          0.005857
2726                                3          0.005857
2882                                4          0.005857
2907                                5          0.005857

```

```

[ ]: # Función para calcular la distancia euclidiana
def haversine(lat1, lon1, lat2, lon2):
    R = 6371 # radio de la Tierra en kilómetros
    phi1 = np.radians(lat1)
    phi2 = np.radians(lat2)
    delta_phi = np.radians(lat2 - lat1)
    delta_lambda = np.radians(lon2 - lon1)
    a = np.sin(delta_phi / 2) ** 2 + np.cos(phi1) * np.cos(phi2) * np.
↳sin(delta_lambda / 2) ** 2
    return R * 2 * np.arctan2(np.sqrt(a), np.sqrt(1 - a))

# Calcular la distancia entre el hogar del cliente y el comercio
df['Distance_home_merch'] = haversine(df['lat'], df['long'], df['merch_lat'],
↳df['merch_long'])

# Calcular la distancia entre el comercio de la última transacción y el
↳comercio actual
df = df.sort_values(by=['cc_num', 'trans_date_trans_time'])
df['Distance_last_trans_merch'] = df.groupby('cc_num').apply(lambda x:
↳haversine(x['merch_lat'], x['merch_long'], x['merch_lat'].shift(),
↳x['merch_long'].shift()))
.reset_index(level=0, drop=True)

```

```

df['Distance_last_trans_merch'].fillna(0, inplace=True)

# Calcular la distancia promedio y máxima entre el hogar del cliente y los
↳ comercios
df['Avg_distance_home_merch'] = df.groupby('cc_num')['Distance_home_merch'].
↳ transform('mean')
df['Max_distance_home_merch'] = df.groupby('cc_num')['Distance_home_merch'].
↳ transform('max')

# Verificación de las nuevas características
df_geograficas_features = df[['Distance_home_merch',
↳ 'Distance_last_trans_merch', 'Avg_distance_home_merch',
↳ 'Max_distance_home_merch']]
df_geograficas_features.head()

```

```

[ ]:      Distance_home_merch  Distance_last_trans_merch  Avg_distance_home_merch \
1017          127.606239              0.000000          73.500496
2724          110.308921          224.769219          73.500496
2726           21.787261          105.220439          73.500496
2882           87.204215           88.152283          73.500496
2907           74.212965          132.876773          73.500496

      Max_distance_home_merch
1017          137.124827
2724          137.124827
2726          137.124827
2882          137.124827
2907          137.124827

```

```

[ ]: # Calcular la edad del cliente
df['dob'] = pd.to_datetime(df['dob'])
df['Age'] = df['trans_date_trans_time'].dt.year - df['dob'].dt.year

# Crear variables demográficas
df['Is_senior'] = (df['Age'] >= 65).astype(int)
df['Is_teenager'] = ((df['Age'] >= 13) & (df['Age'] <= 19)).astype(int)

# Verificación de las nuevas características
df_demograficas_features = df[['Age', 'Is_senior', 'Is_teenager']]
df_demograficas_features.head()

```

```

[ ]:      Age  Is_senior  Is_teenager
1017   33         0         0
2724   33         0         0
2726   33         0         0
2882   33         0         0
2907   33         0         0

```

```
[ ]: # Variables históricas
df['Total_transactions'] = df.groupby('cc_num')['amt'].transform('count')
df['Total_amount_spent'] = df.groupby('cc_num')['amt'].transform('sum')
df['Avg_transaction_amount_category'] = df.groupby(['cc_num', 'category'])['amt'].transform('mean')
df['Std_transaction_amount_category'] = df.groupby(['cc_num', 'category'])['amt'].transform('std').fillna(0)

# Verificación de las nuevas características
df_historicas_features = df[['Total_transactions', 'Total_amount_spent', 'Avg_transaction_amount_category', 'Std_transaction_amount_category']]
df_historicas_features.head()
```

```
[ ]:      Total_transactions  Total_amount_spent  Avg_transaction_amount_category \
1017                2196          130130.12          45.596296
2724                2196          130130.12          59.779429
2726                2196          130130.12          59.779429
2882                2196          130130.12          56.438434
2907                2196          130130.12          55.924559

      Std_transaction_amount_category
1017                113.527496
2724                15.758267
2726                15.758267
2882                65.433283
2907                53.243029
```

```
[ ]: # Variables de ratio
df['Transaction_amount_to_avg_ratio'] = df['amt'] / df['Avg_transaction_amount']
df['Transaction_amount_to_max_ratio'] = df['amt'] / df['Max_transaction_amount']
df['Transaction_amount_to_min_ratio'] = df['amt'] / df['Min_transaction_amount']

# Verificación de las nuevas características
df_ratio_features = df[['Transaction_amount_to_avg_ratio', 'Transaction_amount_to_max_ratio', 'Transaction_amount_to_min_ratio']]
df_ratio_features.head()
```

```
[ ]:      Transaction_amount_to_avg_ratio  Transaction_amount_to_max_ratio \
1017                0.122684                0.002115
2724                0.893385                0.015401
2726                1.385134                0.023878
2882                0.587096                0.010121
2907                0.458674                0.007907

      Transaction_amount_to_min_ratio
1017                7.127451
2724               51.901961
```


2726	80.470588
2882	34.107843
2907	26.647059

```
[ ]: # Variables de comportamiento
df['Is_first_transaction'] = (df.groupby('cc_num')['trans_date_trans_time'].
    ↪rank(method='first') == 1).astype(int)
df['Is_new_merchant'] = (~df.groupby('cc_num')['merchant'].transform(lambda x:
    ↪x.duplicated(keep='first'))).astype(int)
df['Is_new_category'] = (~df.groupby('cc_num')['category'].transform(lambda x:
    ↪x.duplicated(keep='first'))).astype(int)

# Verificación de las nuevas características
df_comportamiento_features = df[['Is_first_transaction', 'Is_new_merchant',
    ↪'Is_new_category']]
df_comportamiento_features.head()
```

```
[ ]:      Is_first_transaction  Is_new_merchant  Is_new_category
1017                      1                1                1
2724                      0                1                1
2726                      0                1                0
2882                      0                1                1
2907                      0                1                1
```

```
[ ]: # Variables de interacción
df['Hour_amount_interaction'] = df['Transaction_hour'] * df['amt']
df['Day_amount_interaction'] = df['Transaction_day'] * df['amt']
df['Month_amount_interaction'] = df['Transaction_month'] * df['amt']

# Verificación de las nuevas características
df_interaccion_features = df[['Hour_amount_interaction',
    ↪'Day_amount_interaction', 'Month_amount_interaction']]
df_interaccion_features.head()
```

```
[ ]:      Hour_amount_interaction  Day_amount_interaction  \
1017                      87.24                7.27
2724                     423.52             105.88
2726                     656.64             164.16
2882                     417.48              69.58
2907                     353.34              54.36

      Month_amount_interaction
1017                      7.27
2724                     52.94
2726                     82.08
2882                     34.79
2907                     27.18
```

```
[ ]: # Variables de comercio
df['Num_unique_merchants'] = df.groupby('cc_num')['merchant'].
    ↪transform('nunique')
df['Num_unique_categories'] = df.groupby('cc_num')['category'].
    ↪transform('nunique')
df['Avg_transactions_per_merchant'] = df.groupby('merchant')['amt'].
    ↪transform('count') / df.groupby('merchant')['cc_num'].transform('nunique')

# Verificación de las nuevas características
df_comercio_features = df[['Num_unique_merchants', 'Num_unique_categories',
    ↪'Avg_transactions_per_merchant']]
df_comercio_features.head()
```

```
[ ]:      Num_unique_merchants  Num_unique_categories  \
1017                      642                      14
2724                      642                      14
2726                      642                      14
2882                      642                      14
2907                      642                      14

      Avg_transactions_per_merchant
1017                      2.614826
2724                      4.645399
2726                      4.614144
2882                      3.828947
2907                      4.248175
```

```
[ ]: # Función para determinar si una fecha es festiva
# Función para determinar si una fecha es festiva (solo mes y día)
def is_holiday(date):
    holidays = [('01-01', '12-25')] # Ejemplo de días festivos (mes-día)
    return date.strftime('%m-%d') in holidays

# Variables de contexto
df['Is_holiday'] = df['trans_date_trans_time'].apply(is_holiday).astype(int)
df['Is_weekend_transaction'] = df['Transaction_day_of_week'].isin([5, 6]).
    ↪astype(int)
df['Merchant_transaction_count'] = df.groupby('merchant')['amt'].
    ↪transform('count')
df['Category_transaction_count'] = df.groupby('category')['amt'].
    ↪transform('count')
df['Is_high_risk_merchant'] = df['merchant'].isin(['fraud_Rippin, Kub and
    ↪Mann', 'fraud_Stokes LLC and Sons']).astype(int) # Ejemplo de comercios de
    ↪alto riesgo

# Verificación de las nuevas características
```

```
df_contexto_features = df[['Is_holiday', 'Is_weekend_transaction',
    ↪ 'Merchant_transaction_count', 'Category_transaction_count',
    ↪ 'Is_high_risk_merchant']]
df_contexto_features.head()
```

```
[ ]:      Is_holiday  Is_weekend_transaction  Merchant_transaction_count  \
1017           0              0              1799
2724           0              0              3786
2726           0              0              3719
2882           0              0              3201
2907           0              0              3492
```

```
      Category_transaction_count  Is_high_risk_merchant
1017              90654              0
2724             188029              0
2726             188029              0
2882             161727              0
2907             175460              0
```

```
[ ]: # Calcular la tendencia del monto de transacción y la frecuencia de transacción
df = df.sort_values(by=['cc_num', 'trans_date_trans_time'])
df['Transaction_amount_trend'] = df.groupby('cc_num')['amt'].transform(lambda x:
    ↪ x.rolling(window=5, min_periods=1).mean())
df['Transaction_frequency_trend'] = df.
    ↪ groupby('cc_num')['trans_date_trans_time'].transform(lambda x: x.diff().dt.
    ↪ total_seconds().rolling(window=5, min_periods=1).mean())

# Calcular el tiempo desde la última transacción fraudulenta del mismo cliente
df['Time_since_last_fraud'] = df.groupby('cc_num')['trans_date_trans_time'].
    ↪ transform(lambda x: x.diff().fillna(pd.Timedelta(seconds=0)).dt.
    ↪ total_seconds())

# Calcular el ratio de transacciones fraudulentas del cliente al total de
    ↪ transacciones
df['Fraud_ratio'] = df.groupby('cc_num')['is_fraud'].transform('mean')

# Calcular el intervalo promedio y la desviación estándar del intervalo entre
    ↪ transacciones del cliente
df['Avg_transaction_interval'] = df.groupby('cc_num')['trans_date_trans_time'].
    ↪ transform(lambda x: x.diff().mean().total_seconds())
df['Std_transaction_interval'] = df.groupby('cc_num')['trans_date_trans_time'].
    ↪ transform(lambda x: x.diff().std().total_seconds())

# Indicador de si el monto de la transacción es significativamente mayor que el
    ↪ promedio del cliente
df['Is_large_amount'] = (df['amt'] > df['Avg_transaction_amount']).astype(int)
```

```

# Indicador de actividad fraudulenta reciente
df['Recent_fraud_activity'] = df.groupby('cc_num')['is_fraud'].transform(lambda
    ↪x: x.rolling(window=5, min_periods=1).sum())

# Indicador de si la transacción ocurre durante horas de alto riesgo (ejemplo:
    ↪0-6 AM)
df['High_risk_time'] = df['Transaction_hour'].isin([0, 1, 2, 3, 4, 5, 6]).
    ↪astype(int)

# Verificación de las nuevas características
df_avanzadas_features = df[['Transaction_amount_trend',
    ↪'Transaction_frequency_trend', 'Time_since_last_fraud',
    ↪'Fraud_ratio', 'Avg_transaction_interval',
    ↪'Std_transaction_interval',
    ↪'Is_large_amount', 'Recent_fraud_activity',
    ↪'High_risk_time']]
df_avanzadas_features.head()

```

```

[ ]:      Transaction_amount_trend  Transaction_frequency_trend  \
1017                7.270                      NaN
2724               30.105                71862.000000
2726               47.430                36010.500000
2882               44.270                28619.666667
2907               40.852                21952.750000

      Time_since_last_fraud  Fraud_ratio  Avg_transaction_interval  \
1017                 0.0      0.004098      28722.084282
2724              71862.0      0.004098      28722.084282
2726               159.0      0.004098      28722.084282
2882             13838.0      0.004098      28722.084282
2907              1952.0      0.004098      28722.084282

      Std_transaction_interval  Is_large_amount  Recent_fraud_activity  \
1017          32251.693854              0              0.0
2724          32251.693854              0              0.0
2726          32251.693854              1              0.0
2882          32251.693854              0              0.0
2907          32251.693854              0              0.0

      High_risk_time
1017              0
2724              0
2726              0
2882              0
2907              0

```

0.2.3 Análisis Exploratorio de Datos

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Exploración inicial del dataset
df.head()
df.info()
df.describe()

# Verificar valores nulos
df.isnull().sum()


# Visualización de la distribución de las clases
sns.countplot(x='is_fraud', data=df)
plt.title('Distribución de Transacciones (Normales vs Fraudulentas)')
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1852394 entries, 1017 to 1850558
Data columns (total 71 columns):
```


#	Column	Dtype
0	amt	float64
1	lat	float64
2	long	float64
3	city_pop	float64
4	unix_time	float64
5	merch_lat	float64
6	merch_long	float64
7	is_fraud	int64
8	amt_month	float64
9	amt_year	float64
10	amt_month_shopping_net_spend	float64
11	count_month_shopping_net	float64
12	first_time_at_merchant	float64
13	Transaction_hour	float64
14	Transaction_day	float64
15	Transaction_month	float64
16	Transaction_year	int32
17	Transaction_day_of_week	float64
18	Transaction_weekend	float64
19	Transaction_quarter	float64
20	Days_since_last_transaction	float64
21	Days_since_first_transaction	float64
22	Transaction_amount_log	float64

23	Transaction_amount_square	float64
24	Transaction_amount_sqrt	float64
25	Transaction_amount_cubed	float64
26	Avg_transaction_amount	float64
27	Std_transaction_amount	float64
28	Min_transaction_amount	float64
29	Max_transaction_amount	float64
30	Num_transactions_last_day	float64
31	Num_transactions_last_week	float64
32	Num_transactions_last_month	float64
33	Avg_transactions_per_day	float64
34	Distance_home_merch	float64
35	Distance_last_trans_merch	float64
36	Avg_distance_home_merch	float64
37	Max_distance_home_merch	float64
38	Age	float64
39	Is_senior	float64
40	Is_teenager	float64
41	Total_transactions	float64
42	Total_amount_spent	float64
43	Avg_transaction_amount_category	float64
44	Std_transaction_amount_category	float64
45	Transaction_amount_to_avg_ratio	float64
46	Transaction_amount_to_max_ratio	float64
47	Transaction_amount_to_min_ratio	float64
48	Is_first_transaction	float64
49	Is_new_merchant	float64
50	Is_new_category	float64
51	Hour_amount_interaction	float64
52	Day_amount_interaction	float64
53	Month_amount_interaction	float64
54	Num_unique_merchants	float64
55	Num_unique_categories	float64
56	Avg_transactions_per_merchant	float64
57	Is_holiday	float64
58	Is_weekend_transaction	float64
59	Merchant_transaction_count	float64
60	Category_transaction_count	float64
61	Is_high_risk_merchant	float64
62	Transaction_amount_trend	float64
63	Transaction_frequency_trend	float64
64	Time_since_last_fraud	float64
65	Fraud_ratio	float64
66	Avg_transaction_interval	float64
67	Std_transaction_interval	float64
68	Is_large_amount	float64
69	Recent_fraud_activity	float64
70	High_risk_time	float64

dtypes: float64(69), int32(1), int64(1)
memory usage: 1010.5 MB



```
[ ]: # Matriz de correlación
numerical_columns = df.select_dtypes(include=[float, int]).columns
corr_matrix = df[numerical_columns].corr()
plt.figure(figsize=(15, 10))
sns.heatmap(corr_matrix, annot=True, fmt='.2f')
plt.title('Matriz de Correlación')
plt.show()
```





```
[ ]: # Gráfico de violín para mostrar la distribución de montos por año
plt.figure(figsize=(14, 7))
sns.violinplot(x='Transaction_year', y='amt', hue='is_fraud', data=df,
               ↪split=True)
plt.title('Distribución de Monto de Transacciones por Año')
plt.show()


# Distribución de montos de transacciones por día de la semana
plt.figure(figsize=(14, 7))
sns.boxplot(x='Transaction_day_of_week', y='amt', hue='is_fraud', data=df)
plt.title('Distribución de Montos de Transacciones por Día de la Semana')
plt.show()

# Histograma de las cantidades de transacciones
plt.figure(figsize=(14, 7))
sns.histplot(df[df['is_fraud'] == 0]['amt'], bins=50, kde=True, color='blue',
             ↪label='No Fraud')
sns.histplot(df[df['is_fraud'] == 1]['amt'], bins=50, kde=True, color='red',
             ↪label='Fraud')
```



```
plt.title('Histograma de Cantidades de Transacciones')
plt.legend()
plt.show()
```





0.2.4 Estandarizacion/Normalizacion de los Datos

```
[ ]: from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE

# Llenar valores nulos
for column in df.columns:
    if df[column].dtype == 'object':
        df[column].fillna(df[column].mode()[0], inplace=True) # Llenar con el
        ↪ valor más frecuente para categóricas
    else:
        df[column].fillna(df[column].median(), inplace=True) # Llenar con la
        ↪ mediana para numéricas

# Eliminar variables no necesarias
df = df.drop(columns=['cc_num', 'trans_date_trans_time', 'merchant',
        ↪ 'category', 'gender', 'first', 'last', 'street', 'city', 'state', 'zip',
        ↪ 'job', 'dob', 'trans_num'])

# Seleccionar características a escalar (excluyendo la columna 'is_fraud' y
        ↪ 'year')
features_to_scale = df.columns.difference(['is_fraud', 'Transaction_year'])

# Inicializar el escalador
scaler = StandardScaler()

# Aplicar el escalador a las características seleccionadas
```

```

df[features_to_scale] = scaler.fit_transform(df[features_to_scale])

# Verificación de los datos normalizados/estandarizados
df[features_to_scale].head()

# Dividir el dataset en características (X) y la variable objetivo (y)
X = df.drop(columns=['is_fraud'])
y = df['is_fraud']

# Asegurarse de que no hay NaNs
df = df.fillna(0)

# Aplicar SMOTE para balancear las clases
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

```

0.2.5 Preparando el Dataset de Entrenamiento

```

[ ]: from sklearn.model_selection import train_test_split

# Dividir el dataset en Train, Dev y Test
X_train, X_temp, y_train, y_temp = train_test_split(X_resampled, y_resampled,
    ↪test_size=0.4, random_state=42)
X_dev, X_test, y_dev, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
    ↪random_state=42)

# Convertir los datos a float (ya deberían estar limpios de NaNs)
X_train = X_train.astype(float)
X_dev = X_dev.astype(float)
X_test = X_test.astype(float)

y_train = y_train.astype(int)
y_dev = y_dev.astype(int)
y_test = y_test.astype(int)

print(f'Tamaño del conjunto de entrenamiento: {X_train.shape[0]}')
print(f'Tamaño del conjunto de validación: {X_dev.shape[0]}')
print(f'Tamaño del conjunto de prueba: {X_test.shape[0]}')

```

```

Tamaño del conjunto de entrenamiento: 2211291
Tamaño del conjunto de validación: 737097
Tamaño del conjunto de prueba: 737098

```

0.2.6 Entrenamiento Normal de los Modelos

```
[ ]: !pip install xgboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-  
packages (2.0.3)  
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages  
(from xgboost) (1.25.2)  
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages  
(from xgboost) (1.11.4)
```

```
[ ]: !nvidia-smi
```

```
Fri May 24 05:57:30 2024
```

```
+-----+  
-----+  
| NVIDIA-SMI 535.104.05                Driver Version: 535.104.05   CUDA Version:  
12.2          |  
+-----+-----+-----+  
-----+  
| GPU   Name                               Persistence-M | Bus-Id        Disp.A | Volatile  
Uncorr. ECC |  
| Fan   Temp   Perf           Pwr:Usage/Cap |      Memory-Usage | GPU-Util  
Compute M. |  
|               |                          |                      |  
MIG M. |  
+=====+  
=====+  
|    0   NVIDIA L4                               Off | 00000000:00:03.0 Off |  
0 |  
| N/A    64C    P8              14W / 72W |      1MiB / 23034MiB |      0%  
Default |  
|               |                          |                      |  
N/A |  
+-----+-----+-----+  
-----+  
  
+-----+  
-----+  
| Processes:  
|  
| GPU   GI    CI          PID    Type    Process name                        GPU  
Memory |  
|       ID    ID                                   |  
Usage   |  
+=====+  
=====+  
| No running processes found
```

```
|
+-----+
-----+
```

```
[ ]: # This get the RAPIDS-Colab install files and test check your GPU. Run this
      ↪ and the next cell only.
      # Please read the output of this cell. If your Colab Instance is not RAPIDS
      ↪ compatible, it will warn you and give you remediation steps.
      !git clone https://github.com/rapidsai/rapidsai-csp-utils.git
      !python rapidsai-csp-utils/colab/pip-install.py
```

```
Cloning into 'rapidsai-csp-utils'...
remote: Enumerating objects: 481, done.
remote: Counting objects: 100% (212/212), done.
remote: Compressing objects: 100% (121/121), done.
remote: Total 481 (delta 143), reused 124 (delta 91), pack-reused 269
Receiving objects: 100% (481/481), 133.58 KiB | 759.00 KiB/s, done.
Resolving deltas: 100% (245/245), done.
Collecting pynvml
  Downloading pynvml-11.5.0-py3-none-any.whl (53 kB)
    53.1/53.1 kB 2.2 MB/s eta 0:00:00
Installing collected packages: pynvml
Successfully installed pynvml-11.5.0
*****
Woo! Your instance has a NVIDIA L4 GPU!
We will install the latest stable RAPIDS via pip 24.4.*! Please stand by,
should be quick...
*****

Looking in indexes: https://pypi.org/simple, https://pypi.nvidia.com
Collecting cuml-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cuml-
cuml-cu12-24.4.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(1200.7 MB)
    1.2/1.2 GB 1.7 MB/s eta 0:00:00
Collecting cugraph-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cugraph-cu12/cugraph_cu12-24.4.0-cp310-cp3
10-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1429.1 MB)
    1.4/1.4 GB 1.5 MB/s eta 0:00:00
Collecting cuspatial-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cuspatial-cu12/cuspatial_cu12-24.4.0-cp310
-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (137.8 MB)
    137.8/137.8 MB 6.7 MB/s eta
0:00:00
Collecting cuproj-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cuproj-cu12/cuproj_cu12-24.4.0-cp310-cp310
-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (920 kB)
    920.9/920.9 kB 73.6 MB/s eta
```

```

0:00:00
Collecting cuxfilter-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cuxfilter-
cu12/cuxfilter_cu12-24.4.1-py3-none-any.whl (83 kB)
      83.5/83.5 kB 11.0 MB/s eta 0:00:00
Collecting cucim-cu12==24.4.*
  Downloading https://pypi.nvidia.com/cucim-cu12/cucim_cu12-24.4.0-cp310-cp310-m
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (5.8 MB)
      5.8/5.8 MB 102.2 MB/s eta 0:00:00
Collecting pylibraft-cu12==24.4.*
  Downloading https://pypi.nvidia.com/pylibraft-cu12/pylibraft_cu12-24.4.0-cp310
-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (823.0 MB)
      823.0/823.0 MB 2.2 MB/s eta
0:00:00
Collecting raft-dask-cu12==24.4.*
  Downloading https://pypi.nvidia.com/raft-dask-cu12/raft_dask_cu12-24.4.0-cp310
-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (170.1 MB)
      170.1/170.1 MB 9.3 MB/s eta
0:00:00
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-
packages (3.9.5)
Requirement already satisfied: cudf-cu12==24.4.* in
/usr/local/lib/python3.10/dist-packages (from cuml-cu12==24.4.*) (24.4.1)
Requirement already satisfied: cupy-cuda12x>=12.0.0 in
/usr/local/lib/python3.10/dist-packages (from cuml-cu12==24.4.*) (12.2.0)
Collecting dask-cuda==24.4.* (from cuml-cu12==24.4.*)
  Downloading dask_cuda-24.4.0-py3-none-any.whl (126 kB)
      126.6/126.6 kB 3.5 MB/s eta
0:00:00
Collecting dask-cudf-cu12==24.4.* (from cuml-cu12==24.4.*)
  Downloading https://pypi.nvidia.com/dask-cudf-
cu12/dask_cudf_cu12-24.4.1-py3-none-any.whl (48 kB)
      48.9/48.9 kB 7.4 MB/s eta 0:00:00
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-
packages (from cuml-cu12==24.4.*) (1.4.2)
Requirement already satisfied: numba>=0.57 in /usr/local/lib/python3.10/dist-
packages (from cuml-cu12==24.4.*) (0.58.1)
Collecting rapids-dask-dependency==24.4.* (from cuml-cu12==24.4.*)
  Downloading https://pypi.nvidia.com/rapids-dask-
dependency/rapids_dask_dependency-24.4.1-py3-none-any.whl (15 kB)
Requirement already satisfied: rmm-cu12==24.4.* in
/usr/local/lib/python3.10/dist-packages (from cuml-cu12==24.4.*) (24.4.0)
Requirement already satisfied: scipy>=1.8.0 in /usr/local/lib/python3.10/dist-
packages (from cuml-cu12==24.4.*) (1.11.4)
Collecting treelite==4.1.2 (from cuml-cu12==24.4.*)
  Downloading treelite-4.1.2-py3-none-manylinux2014_x86_64.whl (810 kB)
      810.9/810.9 kB 9.6 MB/s eta
0:00:00

```

Requirement already satisfied: fsspec[http]>=0.6.0 in
 /usr/local/lib/python3.10/dist-packages (from cugraph-cu12==24.4.*) (2023.6.0)

Requirement already satisfied: numpy<2.0a0,>=1.23 in
 /usr/local/lib/python3.10/dist-packages (from cugraph-cu12==24.4.*) (1.25.2)

Collecting pylibcugraph-cu12==24.4.* (from cugraph-cu12==24.4.*)
 Downloading https://pypi.nvidia.com/pylibcugraph-cu12/pylibcugraph_cu12-24.4.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1430.2 MB)
 1.4/1.4 GB 1.5 MB/s eta 0:00:00

Collecting ucx-py-cu12==0.37.* (from cugraph-cu12==24.4.*)
 Downloading https://pypi.nvidia.com/ucx-py-cu12/ucx_py_cu12-0.37.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (7.7 MB)
 7.7/7.7 MB 99.5 MB/s eta 0:00:00

Requirement already satisfied: geopandas>=0.11.0 in
 /usr/local/lib/python3.10/dist-packages (from cuspatial-cu12==24.4.*) (0.13.2)

Requirement already satisfied: bokeh>=3.1 in /usr/local/lib/python3.10/dist-packages (from cuxfilter-cu12==24.4.*) (3.3.4)

Collecting datashader>=0.15 (from cuxfilter-cu12==24.4.*)
 Downloading datashader-0.16.1-py2.py3-none-any.whl (18.3 MB)
 18.3/18.3 MB 54.9 MB/s eta 0:00:00

Requirement already satisfied: holoviews>=1.16.0 in
 /usr/local/lib/python3.10/dist-packages (from cuxfilter-cu12==24.4.*) (1.17.1)

Collecting jupyter-server-proxy (from cuxfilter-cu12==24.4.*)
 Downloading jupyter_server_proxy-4.1.2-py3-none-any.whl (34 kB)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from cuxfilter-cu12==24.4.*) (24.0)

Requirement already satisfied: panel>=1.0 in /usr/local/lib/python3.10/dist-packages (from cuxfilter-cu12==24.4.*) (1.3.8)

Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from cucim-cu12==24.4.*) (8.1.7)

Requirement already satisfied: lazy-loader>=0.1 in
 /usr/local/lib/python3.10/dist-packages (from cucim-cu12==24.4.*) (0.4)

Requirement already satisfied: scikit-image<0.23.0a0,>=0.19.0 in
 /usr/local/lib/python3.10/dist-packages (from cucim-cu12==24.4.*) (0.19.3)

Requirement already satisfied: cuda-python<13.0a0,>=12.0 in
 /usr/local/lib/python3.10/dist-packages (from pylibraft-cu12==24.4.*) (12.2.1)

Requirement already satisfied: cachetools in /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (5.3.3)

Requirement already satisfied: nvtx>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (0.2.10)

Requirement already satisfied: pandas<2.2.2dev0,>=2.0 in
 /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (2.0.3)

Requirement already satisfied: protobuf<5,>=3.20 in
 /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (3.20.3)

Requirement already satisfied: pynvjitlink-cu12 in
 /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (0.2.3)

Requirement already satisfied: pyarrow<15.0.0a0,>=14.0.1 in /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (14.0.2)

Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (13.7.1)

Requirement already satisfied: typing_extensions>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from cudf-cu12==24.4.*->cuml-cu12==24.4.*) (4.11.0)

Collecting pynvml<11.5,>=11.0.0 (from dask-cuda==24.4.*->cuml-cu12==24.4.*)
 Downloading pynvml-11.4.1-py3-none-any.whl (46 kB)
 47.0/47.0 kB 7.4 MB/s eta 0:00:00

Requirement already satisfied: zict>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from dask-cuda==24.4.*->cuml-cu12==24.4.*) (3.0.0)

Collecting dask==2024.1.1 (from rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*)
 Downloading dask-2024.1.1-py3-none-any.whl (1.2 MB)
 1.2/1.2 MB 68.1 MB/s eta 0:00:00

Collecting distributed==2024.1.1 (from rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*)
 Downloading distributed-2024.1.1-py3-none-any.whl (1.0 MB)
 1.0/1.0 MB 74.2 MB/s eta 0:00:00

Collecting dask-expr==0.4.0 (from rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*)
 Downloading dask_expr-0.4.0-py3-none-any.whl (161 kB)
 161.7/161.7 kB 24.1 MB/s eta 0:00:00

Requirement already satisfied: cloudpickle>=1.5.0 in /usr/local/lib/python3.10/dist-packages (from dask==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (2.2.1)

Requirement already satisfied: partd>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from dask==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (1.4.2)

Requirement already satisfied: pyyaml>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from dask==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (6.0.1)

Requirement already satisfied: toolz>=0.10.0 in /usr/local/lib/python3.10/dist-packages (from dask==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (0.12.1)

Requirement already satisfied: importlib-metadata>=4.13.0 in /usr/local/lib/python3.10/dist-packages (from dask==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (7.1.0)

Requirement already satisfied: Jinja2>=2.10.3 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (3.1.4)

Requirement already satisfied: Locket>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (1.0.0)

Requirement already satisfied: msgpack>=1.0.0 in /usr/local/lib/python3.10/dist-

packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (1.0.8)
 Requirement already satisfied: psutil>=5.7.2 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (5.9.5)
 Requirement already satisfied: sortedcontainers>=2.0.5 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (2.4.0)
 Requirement already satisfied: tblib>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (3.0.0)
 Requirement already satisfied: tornado>=6.0.4 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (6.3.3)
 Requirement already satisfied: urllib3>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (2.0.7)
 Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.3.1)
 Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (23.2.0)
 Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.4.1)
 Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (6.0.5)
 Requirement already satisfied: yarll<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (1.9.4)
 Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp) (4.0.3)
 Requirement already satisfied: contourpy>=1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=3.1->cuxfilter-cu12==24.4.*) (1.2.1)
 Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh>=3.1->cuxfilter-cu12==24.4.*) (9.4.0)
 Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=3.1->cuxfilter-cu12==24.4.*) (2024.4.0)
 Requirement already satisfied: cython in /usr/local/lib/python3.10/dist-packages (from cuda-python<13.0a0,>=12.0->pylibraft-cu12==24.4.*) (3.0.10)
 Requirement already satisfied: fastrlock>=0.5 in /usr/local/lib/python3.10/dist-packages (from cupy-cuda12x>=12.0.0->cuml-cu12==24.4.*) (0.8.2)
 Requirement already satisfied: colorcet in /usr/local/lib/python3.10/dist-packages (from datashader>=0.15->cuxfilter-cu12==24.4.*) (3.1.0)
 Requirement already satisfied: multipledispatch in /usr/local/lib/python3.10/dist-packages (from datashader>=0.15->cuxfilter-cu12==24.4.*) (1.0.0)
 Requirement already satisfied: param in /usr/local/lib/python3.10/dist-packages (from datashader>=0.15->cuxfilter-cu12==24.4.*) (2.1.0)
 Collecting pyct (from datashader>=0.15->cuxfilter-cu12==24.4.*)

Downloading pyct-0.5.0-py2.py3-none-any.whl (15 kB)
 Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from datashader>=0.15->cuxfilter-cu12==24.4.*) (2.31.0)
 Requirement already satisfied: xarray in /usr/local/lib/python3.10/dist-packages (from datashader>=0.15->cuxfilter-cu12==24.4.*) (2023.7.0)
 Requirement already satisfied: fiona>=1.8.19 in /usr/local/lib/python3.10/dist-packages (from geopandas>=0.11.0->cuspatial-cu12==24.4.*) (1.9.6)
 Requirement already satisfied: pyproj>=3.0.1 in /usr/local/lib/python3.10/dist-packages (from geopandas>=0.11.0->cuspatial-cu12==24.4.*) (3.6.1)
 Requirement already satisfied: shapely>=1.7.1 in /usr/local/lib/python3.10/dist-packages (from geopandas>=0.11.0->cuspatial-cu12==24.4.*) (2.0.4)
 Requirement already satisfied: pyviz-comms>=0.7.4 in /usr/local/lib/python3.10/dist-packages (from holoviews>=1.16.0->cuxfilter-cu12==24.4.*) (3.0.2)
 Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba>=0.57->cuml-cu12==24.4.*) (0.41.1)
 Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (3.6)
 Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (3.0.0)
 Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (2.0.3)
 Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (0.4.1)
 Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (4.66.4)
 Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=1.0->cuxfilter-cu12==24.4.*) (6.1.0)
 Requirement already satisfied: networkx>=2.2 in /usr/local/lib/python3.10/dist-packages (from scikit-image<0.23.0a0,>=0.19.0->cucim-cu12==24.4.*) (3.3)
 Requirement already satisfied: imageio>=2.4.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image<0.23.0a0,>=0.19.0->cucim-cu12==24.4.*) (2.31.6)
 Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.10/dist-packages (from scikit-image<0.23.0a0,>=0.19.0->cucim-cu12==24.4.*) (2024.5.10)
 Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image<0.23.0a0,>=0.19.0->cucim-cu12==24.4.*) (1.6.0)
 Requirement already satisfied: idna>=2.0 in /usr/local/lib/python3.10/dist-packages (from yarl<2.0,>=1.0->aiohttp) (3.7)
 Requirement already satisfied: jupyter-server>=1.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-server-proxy->cuxfilter-cu12==24.4.*) (1.24.0)
 Collecting simpervisor>=1.0 (from jupyter-server-proxy->cuxfilter-cu12==24.4.*)
 Downloading simpervisor-1.0.0-py3-none-any.whl (8.3 kB)
 Requirement already satisfied: traitlets>=4.2.1 in

/usr/local/lib/python3.10/dist-packages (from jupyter-server-proxy->cuxfilter-cu12==24.4.*) (5.7.1)

Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (from fiona>=1.8.19->geopandas>=0.11.0->cuspatial-cu12==24.4.*) (2024.2.2)

Requirement already satisfied: click-plugins>=1.0 in /usr/local/lib/python3.10/dist-packages (from fiona>=1.8.19->geopandas>=0.11.0->cuspatial-cu12==24.4.*) (1.1.1)

Requirement already satisfied: cligj>=0.5 in /usr/local/lib/python3.10/dist-packages (from fiona>=1.8.19->geopandas>=0.11.0->cuspatial-cu12==24.4.*) (0.7.2)

Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from fiona>=1.8.19->geopandas>=0.11.0->cuspatial-cu12==24.4.*) (1.16.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2>=2.10.3->distributed==2024.1.1->rapids-dask-dependency==24.4.*->cuml-cu12==24.4.*) (2.1.5)

Requirement already satisfied: anyio<4,>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (3.7.1)

Requirement already satisfied: argon2-cffi in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (23.1.0)

Requirement already satisfied: jupyter-client>=6.1.12 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (6.1.12)

Requirement already satisfied: jupyter-core!=5.0.*,>=4.12 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (5.7.2)

Requirement already satisfied: nbconvert>=6.4.4 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (6.5.4)

Requirement already satisfied: nbformat>=5.2.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (5.10.4)

Requirement already satisfied: prometheus-client in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.20.0)

Requirement already satisfied: pyzmq>=17 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (24.0.1)

Requirement already satisfied: Send2Trash in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (1.8.3)

Requirement already satisfied: terminado>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.18.1)

Requirement already satisfied: websocket-client in /usr/local/lib/python3.10/dist-packages (from jupyter-server>=1.0->jupyter-

server-proxy->cuxfilter-cu12==24.4.*) (1.8.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas<2.2.2dev0,>=2.0->cudf-
cu12==24.4.*->cuml-cu12==24.4.*) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas<2.2.2dev0,>=2.0->cudf-cu12==24.4.*->cuml-cu12==24.4.*)
(2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-
packages (from pandas<2.2.2dev0,>=2.0->cudf-cu12==24.4.*->cuml-cu12==24.4.*)
(2024.1)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-
packages (from bleach->panel>=1.0->cuxfilter-cu12==24.4.*) (0.5.1)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-
packages (from linkify-it-py->panel>=1.0->cuxfilter-cu12==24.4.*) (1.0.3)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-
packages (from markdown-it-py->panel>=1.0->cuxfilter-cu12==24.4.*) (0.1.2)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from
requests->datashader>=0.15->cuxfilter-cu12==24.4.*) (3.3.2)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.10/dist-packages (from rich->cudf-cu12==24.4.*->cuml-
cu12==24.4.*) (2.16.1)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.10/dist-
packages (from anyio<4,>=3.1.0->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (1.3.1)
Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-
packages (from anyio<4,>=3.1.0->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (1.2.1)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.10/dist-
packages (from importlib-metadata>=4.13.0->dask==2024.1.1->rapids-dask-
dependency==24.4.*->cuml-cu12==24.4.*) (3.18.2)
Requirement already satisfied: platformdirs>=2.5 in
/usr/local/lib/python3.10/dist-packages (from jupyter-
core!=5.0.*,>=4.12->jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-
cu12==24.4.*) (4.2.2)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages
(from nbconvert>=6.4.4->jupyter-server>=1.0->jupyter-server-proxy->cuxfilter-
cu12==24.4.*) (4.9.4)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-
packages (from nbconvert>=6.4.4->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (4.12.3)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-
packages (from nbconvert>=6.4.4->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (0.7.1)
Requirement already satisfied: entrypoints>=0.2.2 in
/usr/local/lib/python3.10/dist-packages (from nbconvert>=6.4.4->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.4)
Requirement already satisfied: jupyterlab-pygments in

```

/usr/local/lib/python3.10/dist-packages (from nbconvert>=6.4.4->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.3.0)
Requirement already satisfied: mistune<2,>=0.8.1 in
/usr/local/lib/python3.10/dist-packages (from nbconvert>=6.4.4->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.8.4)
Requirement already satisfied: nbclient>=0.5.0 in
/usr/local/lib/python3.10/dist-packages (from nbconvert>=6.4.4->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (0.10.0)
Requirement already satisfied: pandocfilters>=1.4.1 in
/usr/local/lib/python3.10/dist-packages (from nbconvert>=6.4.4->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (1.5.1)
Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-
packages (from nbconvert>=6.4.4->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (1.3.0)
Requirement already satisfied: fastjsonschema>=2.15 in
/usr/local/lib/python3.10/dist-packages (from nbformat>=5.2.0->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (2.19.1)
Requirement already satisfied: jsonschema>=2.6 in
/usr/local/lib/python3.10/dist-packages (from nbformat>=5.2.0->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (4.19.2)
Requirement already satisfied: ptyprocess in /usr/local/lib/python3.10/dist-
packages (from terminado>=0.8.3->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (0.7.0)
Requirement already satisfied: argon2-cffi-bindings in
/usr/local/lib/python3.10/dist-packages (from argon2-cffi->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (21.2.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in
/usr/local/lib/python3.10/dist-packages (from
jsonschema>=2.6->nbformat>=5.2.0->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (2023.12.1)
Requirement already satisfied: referencing>=0.28.4 in
/usr/local/lib/python3.10/dist-packages (from
jsonschema>=2.6->nbformat>=5.2.0->jupyter-server>=1.0->jupyter-server-
proxy->cuxfilter-cu12==24.4.*) (0.35.1)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-
packages (from jsonschema>=2.6->nbformat>=5.2.0->jupyter-server>=1.0->jupyter-
server-proxy->cuxfilter-cu12==24.4.*) (0.18.1)
Requirement already satisfied: cffi>=1.0.1 in /usr/local/lib/python3.10/dist-
packages (from argon2-cffi-bindings->argon2-cffi->jupyter-server>=1.0->jupyter-
server-proxy->cuxfilter-cu12==24.4.*) (1.16.0)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-
packages (from beautifulsoup4->nbconvert>=6.4.4->jupyter-server>=1.0->jupyter-
server-proxy->cuxfilter-cu12==24.4.*) (2.5)
Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-
packages (from cffi>=1.0.1->argon2-cffi-bindings->argon2-cffi->jupyter-
server>=1.0->jupyter-server-proxy->cuxfilter-cu12==24.4.*) (2.22)
Installing collected packages: simpervisor, pynvml, pyct, ucx-py-cu12, treelite,
dask, pylibraft-cu12, distributed, dask-expr, cuproj-cu12, cucim-cu12, rapids-

```

dask-dependency, pylibcugraph-cu12, datashader, cuspatial-cu12, dask-cudf-cu12, dask-cuda, raft-dask-cu12, cuml-cu12, cugraph-cu12, jupyter-server-proxy, cuxfilter-cu12

```
Attempting uninstall: pynvml
  Found existing installation: pynvml 11.5.0
  Uninstalling pynvml-11.5.0:
    Successfully uninstalled pynvml-11.5.0
Attempting uninstall: dask
  Found existing installation: dask 2023.8.1
  Uninstalling dask-2023.8.1:
    Successfully uninstalled dask-2023.8.1
Attempting uninstall: distributed
  Found existing installation: distributed 2023.8.1
  Uninstalling distributed-2023.8.1:
    Successfully uninstalled distributed-2023.8.1
```

Successfully installed cucim-cu12-24.4.0 cugraph-cu12-24.4.0 cuml-cu12-24.4.0 cuproj-cu12-24.4.0 cuspatial-cu12-24.4.0 cuxfilter-cu12-24.4.1 dask-2024.1.1 dask-cuda-24.4.0 dask-cudf-cu12-24.4.1 dask-expr-0.4.0 datashader-0.16.1 distributed-2024.1.1 jupyter-server-proxy-4.1.2 pyct-0.5.0 pylibcugraph-cu12-24.4.0 pylibraft-cu12-24.4.0 pynvml-11.4.1 raft-dask-cu12-24.4.0 rapids-dask-dependency-24.4.1 simpservisor-1.0.0 treelite-4.1.2 ucx-py-cu12-0.37.0

```
*****
The pip install of RAPIDS is complete.
```

Please do not run any further installation from the conda based installation methods, as they may cause issues!

Please ensure that you're pulling from the git repo to remain updated with the latest working install scripts.

Troubleshooting:

- If there is an installation failure, please check back on RAPIDSAI owned templates/notebooks to see how to update your personal files.

- If an installation failure persists when using the latest script, please make an issue on <https://github.com/rapidsai-community/rapidsai-csp-utils>

```
*****
```

```
[ ]: import cuml
from cuml.ensemble import RandomForestClassifier as cuRF
from sklearn.metrics import roc_auc_score, precision_score, recall_score, f1_score, accuracy_score, confusion_matrix, ConfusionMatrixDisplay

# Inicializar el modelo Random Forest de cuML
rf_model = cuRF(
    n_estimators=200,
```

```

    max_depth=10,
    min_samples_split=5,
    min_samples_leaf=2,
    max_features='sqrt',
    random_state=42
)

# Entrenamiento del modelo
rf_model.fit(X_train, y_train)

# Predicción y evaluación del modelo Random Forest
y_pred_rf_proba = rf_model.predict_proba(X_test).to_numpy()[:, 1]
y_pred_rf = rf_model.predict(X_test)

rf_roc_auc = roc_auc_score(y_test, y_pred_rf_proba)
rf_precision = precision_score(y_test, y_pred_rf)
rf_recall = recall_score(y_test, y_pred_rf)
rf_f1 = f1_score(y_test, y_pred_rf)
rf_accuracy = accuracy_score(y_test, y_pred_rf)

rf_metrics = {
    'ROC-AUC': rf_roc_auc,
    'Precision': rf_precision,
    'Recall': rf_recall,
    'F1 Score': rf_f1,
    'Accuracy': rf_accuracy
}

print("Random Forest Metrics:", rf_metrics)


# Matriz de Confusión
rf_cm = confusion_matrix(y_test, y_pred_rf)
ConfusionMatrixDisplay(rf_cm).plot()
plt.title("Random Forest Confusion Matrix")
plt.show()

```

```

Random Forest Metrics: {'ROC-AUC': 0.9999898572277597, 'Precision':
0.999409968540719, 'Recall': 0.9976711235607381, 'F1 Score': 0.9985397890504554,
'Accuracy': 0.998541577917726}

```



```
[ ]: import xgboost as xgb

# Inicializar el modelo XGBoost con GPU
params = {
    'objective': 'binary:logistic',
    'tree_method': 'gpu_hist',
    'eval_metric': 'auc',
    'learning_rate': 0.05,
    'max_depth': 6,
    'min_child_weight': 3,
    'subsample': 0.8,
    'colsample_bytree': 0.8,
    'gamma': 0.1,
    'alpha': 0.1,
    'lambda': 1,
    'random_state': 42
}
xgb_model = xgb.XGBClassifier(**params)

# Entrenamiento del modelo
```



```

xgb_model.fit(X_train, y_train)

# Predicción y evaluación del modelo XGBoost
y_pred_xgb = xgb_model.predict(X_test)
y_pred_xgb_proba = xgb_model.predict_proba(X_test)[:, 1]

xgb_roc_auc = roc_auc_score(y_test, y_pred_xgb_proba)
xgb_precision = precision_score(y_test, y_pred_xgb)
xgb_recall = recall_score(y_test, y_pred_xgb)
xgb_f1 = f1_score(y_test, y_pred_xgb)
xgb_accuracy = accuracy_score(y_test, y_pred_xgb)

xgb_metrics = {
    'ROC-AUC': xgb_roc_auc,
    'Precision': xgb_precision,
    'Recall': xgb_recall,
    'F1 Score': xgb_f1,
    'Accuracy': xgb_accuracy
}

print("XGBoost Metrics:", xgb_metrics)

# Matriz de Confusión
xgb_cm = confusion_matrix(y_test, y_pred_xgb)
ConfusionMatrixDisplay(xgb_cm).plot()
plt.title("XGBoost Confusion Matrix")
plt.show()

```

/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[06:23:38] WARNING: /workspace/src/common/error_msg.cc:27: The tree method
`gpu_hist` is deprecated since 2.0.0. To use GPU training, set the `device`
parameter to CUDA instead.

E.g. tree_method = "hist", device = "cuda"

```
warnings.warn(smsg, UserWarning)
```

/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[06:23:40] WARNING: /workspace/src/common/error_msg.cc:27: The tree method
`gpu_hist` is deprecated since 2.0.0. To use GPU training, set the `device`
parameter to CUDA instead.

E.g. tree_method = "hist", device = "cuda"

```
warnings.warn(smsg, UserWarning)
```

/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[06:23:40] WARNING: /workspace/src/common/error_msg.cc:58: Falling back to
prediction using DMatrix due to mismatched devices. This might lead to higher

memory usage and slower performance. XGBoost is running on: cuda:0, while the input data is on: cpu.

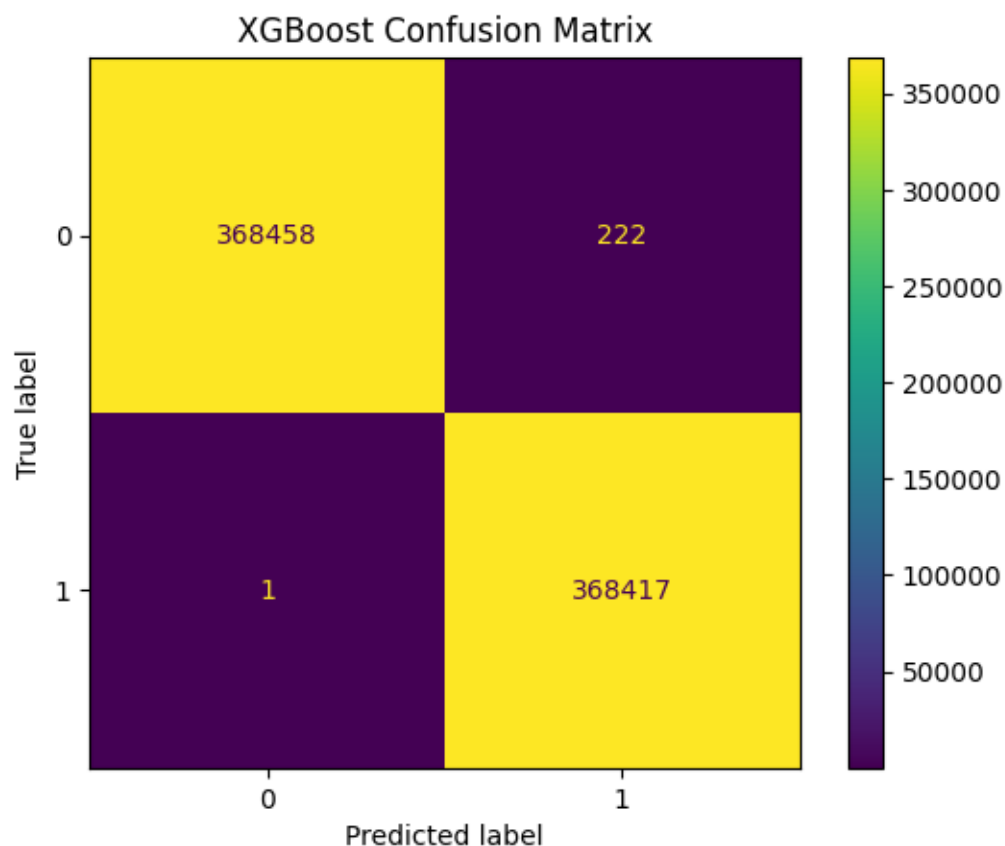
Potential solutions:

- Use a data structure that matches the device ordinal in the booster.
- Set the device for booster before call to inplace_predict.

This warning will only be shown once.

```
warnings.warn(smsg, UserWarning)
```

```
XGBoost Metrics: {'ROC-AUC': 0.9999964878502449, 'Precision':  
0.9993977848247201, 'Recall': 0.9999972856917957, 'F1 Score':  
0.9996974453807507, 'Accuracy': 0.9996974622099096}
```



0.2.7 Entrenamiento Incremental de los Modelos

```
[ ]: # Inicializar el modelo Random Forest de cuML  
rf_model = cuRF(  
    n_estimators=200,  
    max_depth=10,
```

```

        min_samples_split=5,
        min_samples_leaf=2,
        max_features='sqrt',
        random_state=42
    )

    # Entrenamiento incremental manual por año
    years = sorted(X_train['Transaction_year'].unique())
    for year in years:
        # Filtrar los datos por año
        X_train_year = X_train[X_train['Transaction_year'] == year]
        y_train_year = y_train.loc[X_train_year.index]

        # Entrenar el modelo con los datos del año actual
        rf_model.fit(X_train_year.drop(columns=['Transaction_year']), y_train_year)

    # Predicción y evaluación del modelo Random Forest
    y_pred_rf_proba = rf_model.predict_proba(X_test.
        ↪drop(columns=['Transaction_year'])).to_numpy()[:, 1]
    y_pred_rf = rf_model.predict(X_test.drop(columns=['Transaction_year']))

    rf_roc_auc = roc_auc_score(y_test, y_pred_rf_proba)
    rf_precision = precision_score(y_test, y_pred_rf)
    rf_recall = recall_score(y_test, y_pred_rf)
    rf_f1 = f1_score(y_test, y_pred_rf)
    rf_accuracy = accuracy_score(y_test, y_pred_rf)

    rf_metrics = {
        'ROC-AUC': rf_roc_auc,
        'Precision': rf_precision,
        'Recall': rf_recall,
        'F1 Score': rf_f1,
        'Accuracy': rf_accuracy
    }

    print("Incremental Random Forest Metrics:", rf_metrics)

    # Matriz de Confusión
    rf_cm = confusion_matrix(y_test, y_pred_rf)
    ConfusionMatrixDisplay(rf_cm).plot()
    plt.title("Incremental Random Forest Confusion Matrix")
    plt.show()

```

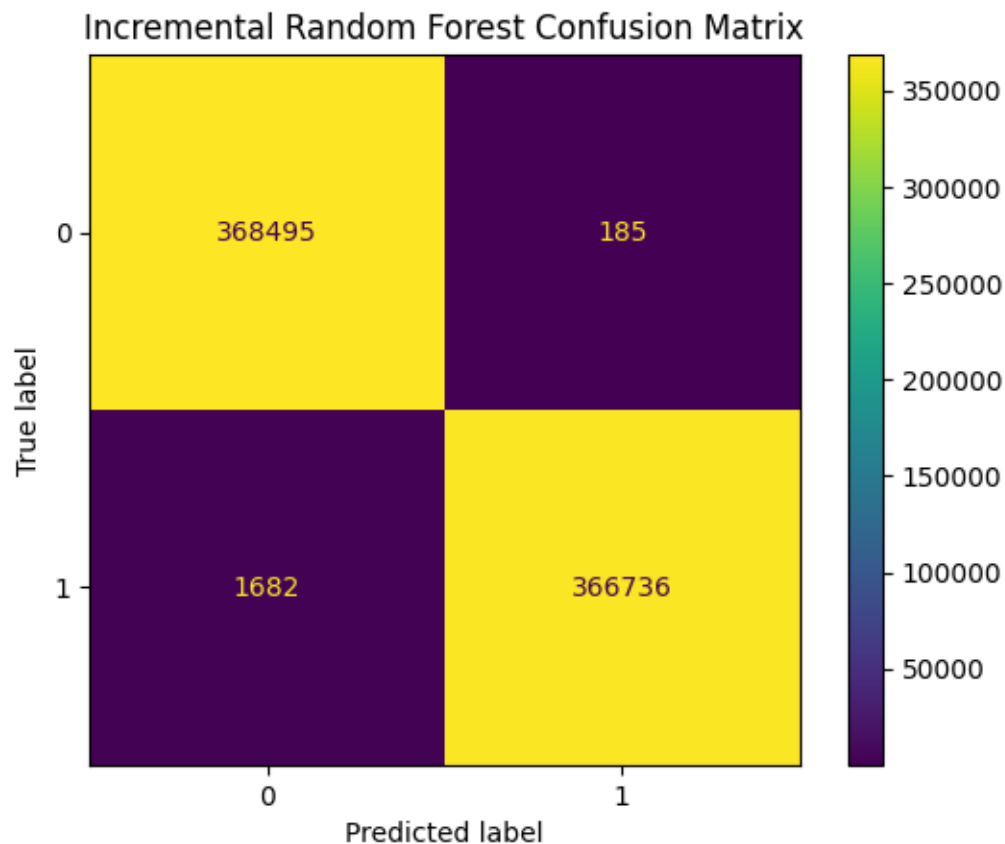
/usr/local/lib/python3.10/dist-packages/cuml/internals/api_decorators.py:344:
 UserWarning: For reproducible results in Random Forest Classifier or for almost
 reproducible results in Random Forest Regressor, n_streams=1 is recommended. If
 n_streams is > 1, results may vary due to stream/thread timing differences, even

```

when random_state is set
    return func(**kwargs)
/usr/local/lib/python3.10/dist-packages/cuml/internals/api_decorators.py:188:
UserWarning: To use pickling first train using float32 data to fit the estimator
    ret = func(*args, **kwargs)
/usr/local/lib/python3.10/dist-packages/cuml/internals/api_decorators.py:188:
UserWarning: To use pickling first train using float32 data to fit the estimator
    ret = func(*args, **kwargs)

Incremental Random Forest Metrics: {'ROC-AUC': 0.9999877081329369, 'Precision':
0.9994958042739446, 'Recall': 0.9954345336004212, 'F1 Score':
0.9974610349784249, 'Accuracy': 0.9974670939278087}

```



```

[ ]: # Inicializar los parámetros del modelo XGBoost
params = {
    'objective': 'binary:logistic',
    'tree_method': 'gpu_hist',
    'eval_metric': 'auc',
    'learning_rate': 0.05,
    'max_depth': 6,
    'min_child_weight': 3,

```

```

        'subsample': 0.8,
        'colsample_bytree': 0.8,
        'gamma': 0.1,
        'alpha': 0.1,
        'lambda': 1,
        'random_state': 42
    }

    # Inicializar el modelo XGBoost
    xgb_model = None

    # Entrenamiento incremental manual por año
    years = sorted(X_train['Transaction_year'].unique())
    for year in years:
        # Filtrar los datos por año
        X_train_year = X_train[X_train['Transaction_year'] == year]
        y_train_year = y_train.loc[X_train_year.index]

        dtrain = xgb.DMatrix(X_train_year.drop(columns=['Transaction_year']),
                               label=y_train_year)

        # Entrenar el modelo con los datos del año actual
        if xgb_model is None:
            xgb_model = xgb.train(params, dtrain, num_boost_round=100)
        else:
            xgb_model = xgb.train(params, dtrain, num_boost_round=100,
                                   xgb_model=xgb_model)

    # Predicción y evaluación del modelo XGBoost
    dtest = xgb.DMatrix(X_test.drop(columns=['Transaction_year']))
    y_pred_xgb_proba = xgb_model.predict(dtest)
    y_pred_xgb = (y_pred_xgb_proba > 0.5).astype(int)

    xgb_roc_auc = roc_auc_score(y_test, y_pred_xgb_proba)
    xgb_precision = precision_score(y_test, y_pred_xgb)
    xgb_recall = recall_score(y_test, y_pred_xgb)
    xgb_f1 = f1_score(y_test, y_pred_xgb)
    xgb_accuracy = accuracy_score(y_test, y_pred_xgb)

    xgb_metrics = {
        'ROC-AUC': xgb_roc_auc,
        'Precision': xgb_precision,
        'Recall': xgb_recall,
        'F1 Score': xgb_f1,
        'Accuracy': xgb_accuracy
    }

```

```
print("XGBoost Metrics:", xgb_metrics)

# Matriz de Confusión
xgb_cm = confusion_matrix(y_test, y_pred_xgb)
ConfusionMatrixDisplay(xgb_cm).plot()
plt.title("XGBoost Confusion Matrix")
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[07:47:14] WARNING: /workspace/src/common/error_msg.cc:27: The tree method
`gpu_hist` is deprecated since 2.0.0. To use GPU training, set the `device`
parameter to CUDA instead.
```

E.g. `tree_method = "hist", device = "cuda"`

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[07:47:18] WARNING: /workspace/src/common/error_msg.cc:27: The tree method
`gpu_hist` is deprecated since 2.0.0. To use GPU training, set the `device`
parameter to CUDA instead.
```

E.g. `tree_method = "hist", device = "cuda"`

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning:
[07:47:20] WARNING: /workspace/src/common/error_msg.cc:27: The tree method
`gpu_hist` is deprecated since 2.0.0. To use GPU training, set the `device`
parameter to CUDA instead.
```

E.g. `tree_method = "hist", device = "cuda"`

```
warnings.warn(smsg, UserWarning)
XGBoost Metrics: {'ROC-AUC': 0.9999994098507341, 'Precision':
0.9998235939077478, 'Recall': 0.999959285376936, 'F1 Score': 0.9998914350387984,
'Accuracy': 0.9998914662636448}
```

